

Nakamoto Consensus

Zeta Avarikioti


georgia.avarikioti@tuwien.ac.at

November 19th, 2025

Overview


- ▶ Nakamoto consensus recap
- ▶ Key observations for Bitcoin
(assumptions, efficiency, decentralization, etc)
- ▶ Security proof (simplified)
- ▶ Security under network delays
- ▶ Selfish mining attack

Nakamoto consensus recap



- ▶ Mine on longest chain, i.e., repeatedly $\text{hash}(\text{tx}, \text{nonce}, \text{previous_hash})$ until hit the target
 - ▶ Upon having a longest chain, send to all (via peer-to-peer network)
 - ▶ Commit blocks buried deep

Why does it work?



- ▶ A unique longest chain keeps growing faster than all other chains
 - ▶ Hence, once a block is buried deep, it is less likely to be “forked off”
 - ▶ More rigorous proof will follow

Key Observations


Interesting questions

- ▶ Under which assumptions is Bitcoin secure?
 - ▶ Timing model
 - ▶ Signatures
- ▶ Does peer-to-peer help or hurt?
 - ▶ Latency
 - ▶ Communication
- ▶ What is the role of PoW?
- ▶ Is Bitcoin permissionless and decentralised?

Model of Bitcoin

- ▶ Timing model
 - ▶ Secure under synchrony
 - ▶ Insecure under partial synchrony or asynchrony
 - ▶ Some model in between?

The CAP Theorem




Model of Bitcoin

- ▶ Timing model
 - ▶ Secure under synchrony
 - ▶ Insecure under partial synchrony or asynchrony
 - ▶ Some model in between?
- ▶ Authentication
 - ▶ Signatures are not necessary for consensus
 - ▶ Yet overcome the $1/3$ fault bound (FLM), why?

[PSL80,FLM85] If $f \geq n/3$, no deterministic protocol for Byzantine broadcast satisfies both agreement and validity in the synchronous model.

Efficiency of Bitcoin

- ▶ Very high latency
 - ▶ For $\exp(-k)$ error probability, wait for k blocks
 - ▶ Blocks are slowly arriving
- ▶ Communication: $O(ndB)$
 - ▶ n : number of nodes
 - ▶ d : number of neighbours
 - ▶ B : block size (4MB)
 - ▶ Without peer-to-peer: $O(n^2B)$
- ▶ Enormous energy consumption



Roles of PoW


- ▶ Sybil resilience
- ▶ Leader election
 - ▶ Unpredictable
 - ▶ Rate limiting
- ▶ Make equivocation hard
 - ▶ Somewhat resembles signatures

Bitcoin: permissionless, decentralized?

- ▶ Permissionless

- ▶ Dynamic network 
- ▶ Anyone can join/no IDs, no permissions 

Bitcoin Mining Pools Distribution Pie Chart



- ▶ Decentralized

- ▶ Operated by many nodes 
- ▶ Nodes are not controlled by only a few 

Security Proof

(a simplified one)

Objectives

- ▶ Security of Bitcoin
 - ▶ **Safety:** A block b is safe if, once it is confirmed it remains confirmed at the view of all honest nodes at all future times, irrespective of the adversary attack.
 - ▶ **Liveness:** A non-zero fraction of honest blocks are confirmed.

Safety is:


- 1) Consistency among honest nodes at any time
- 2) Self-consistency of each honest node across time

General Assumptions

- ▶ (A1) No node has knowledge of the genesis block prior to the deployment of the protocol (trusted setup).
- ▶ (A2) It is easy for all nodes to verify whether a given node is the leader of a given round.
- ▶ (A3) No node can influence the probability with which it is selected as the leader of a round.
- ▶ (A4) Every block must claim as its predecessor some block that belongs to a previous round.
(Only PoW) Each leader produces exactly one block per round.
- ▶ (A5) At all times, all honest nodes know about the exact same set of blocks and predecessors (telepathy).

Adversarial Deviations


- ▶ Empty blocks
- ▶ Delayed block announcements
- ▶ Deliberate forking
- ▶ Combination of the above!
 - ▶ e.g. Private double-spent attack



Safety of Nakamoto


under 0 network delay

k-deep confirmation rule




- ▶ Suppose 1-deep confirmation rule:
 - ▶ Alice pays Bob in block b (tx)
 - ▶ Alice starts immediately mining in private on top of b' (tx')
 - ▶ Bob sends car to Alice
 - ▶ Alice reveals her private chain
 - ▶ Block b is pushed out of the longest chain

k-deep confirmation rule




- ▶ What if miners keep building on block b ?
 - ▶ Healability after network partition
 - ▶ Safety under network delays
 - ▶ Safety for bootstrapping new nodes
- ▶ The attack is feasible even from weak attackers → It would erode people's trust in the system!

k-deep confirmation rule



- ▶ k-deep confirmation rule:
 - ▶ Alice pays Bob in block b
 - ▶ k-1 blocks are build on top of block b
 - ▶ Bob sends car to Alice


k-deep confirmation rule



- ▶ Why k-deep confirmation rule works?
 - ▶ The **probability of tx being pushed out of the longest chain decays exponentially as k grows!**
 - ▶ The bound holds as long as the adversary holds strictly less than half of the total mining power, $\beta < 1/2$;
Else, the adversary can always override the honest chain.


Confirmation error probability

Confirmation
error prob.



Log-linear plot of confirmation error probability as a function of depth k for adversarial hash power 30%

Definitions



- ▶ $B_k(G)$: the longest chain of in-tree G , with the last k blocks removed.
- ▶ w -balanced leader sequence: A sequence of leaders is w -balanced when every window of length at least w , honest leaders outnumber adversarial leaders.

Sketch Proofs

Methodology

- ▶ Random Leader Selection Implies Balanced Sequences (A3 critical)
- ▶ Balanced Leader Sequences Imply Security Properties
 - ▶ Common Prefix (A1, A4 critical)
 - ▶ Chain Quality
 - ▶ Chain Growth

Sketch Proofs


Random leaders imply balanced sequences

- ▶ PR1: Probability a window w is not w -balanced $\leq e^{cw}$
 - ▶ c depends on how close is the adversary to 50%
- ▶ PR2: Probability any window of length w not being w -balanced is the sum of PR1 in all possible windows in a timeframe (union bound)
 - ▶ $\approx T^2$ many windows of at least length w in a timeframe T
 - ▶ $\text{PR2} \leq T^2 e^{cw}$
 - ▶ Hence, the probability all windows are w -balanced is $\geq 1 - T^2 e^{cw}$

Intuition: Law of large numbers

Sketch Proofs


Balanced leader sequences imply common prefix



- ▶ Towards contradiction, assuming $(2k + 2)$ –balanced leader sequence
 - ▶ There are 2 chains that differ for more than k blocks.
 - ▶ Let B_0 the most recent ancestor of B^* that was produced by an honest leader (it always exists due to A1)
 - ▶ We will count leaders and show that there are more Byzantine than honest.

Sketch Proofs


Balanced leader sequences imply common prefix



- ▶ At the same height, we cannot have two honest blocks!
 - ▶ Synchrony (A5)
 - ▶ Honest follow the Nakamoto consensus (i.e., build on the longest chain)

Sketch Proofs

Balanced leader sequences imply common prefix




- ▶ Hence, after B^* , at the same height:
 - ▶ One honest block (H) and at least one adversarial (A) (cf. height 1,2)
 - ▶ No honest block and multiple adversarial (cf. height 3)

$$A \geq H$$

Sketch Proofs

Balanced leader sequences imply common prefix




So, after round r_0 :

1. There have been at least $2k + 2$ leaders

Sketch Proofs

Balanced leader sequences imply common prefix




So, after round r_0 :

1. There have been at least $2k + 2$ leaders
2. At most $h - h_0$ leaders could have been honest

Sketch Proofs

Balanced leader sequences imply common prefix




So, after round r_0 :

1. There have been at least $2k + 2$ leaders
2. At most $h - h_0$ leaders could have been honest
3. At least $h - h_0$ leaders must have been Byzantine
(all before h^* & at least one per higher after h^*)

Sketch Proofs

Balanced leader sequences imply common prefix




So, after round r_0 :

1. There have been at least $2k + 2$ leaders
2. **Byzantine** leaders outnumber the **honest** leaders

Not $(2k + 2)$ –balanced leader sequence (contradiction).

Sketch Proofs

Finality




Common prefix implies that if I crop the last k blocks on my local chain, the remaining **chain is final** (i.e., common among all honest nodes across time).

Liveness

Liveness

- ▶ Recall **liveness** definition: *A non-zero fraction of honest blocks are confirmed.*
- ▶ Observations:
 - ▶ Non honest blocks can be empty or junk
 - ▶ Non zero fraction guarantees that any honest transaction will be eventually included


Chain quality



Chain quality (CQ): The long term average fraction of honest blocks on the longest chain.

Example: $CQ = 4/6$


Chain quality



Chain quality (CQ): The long term average fraction of honest blocks on the longest chain.

Liveness: $CQ > 0$


Chain growth



Chain growth (g): The long term average growth rate of the longest chain (blocks/sec).

- ▶ Let m_a, m_h be the mining rate of adversary and honest (blocks/sec)
- ▶ If the adversary inserts all his blocks in the longest chain, $g \geq m_a$

Chain quality lower bound



- ▶ We have: $CQ \geq \min_g \frac{g - m_a}{g}$
 - ▶ g can be manipulated by the adversary
 - ▶ If the adversary is following the protocol: $g = m_a + m_h$
 - ▶ The minimum growth though depends on the honest mining $g \geq m_h$

$$CQ \geq \min_{g \geq m_h} \frac{g - m_a}{g} = 1 - \frac{m_a}{m_h} \quad \Rightarrow \quad CQ > 0 \text{ if } m_h > m_a$$

Security conclusions

Bitcoin is live if $\alpha < 1/2$

Bitcoin is safe if $\alpha < 1/2$

Bitcoin is secure if $\alpha < 1/2$


Safety under...

bounded network delay

Network delays


- ▶ Bitcoin block creation time: 10min
- ▶ Bitcoin propagation delay: \sim 12sec
- ▶ Δ -synchronous network model: every node receives the data (block) within the bounded delay, denoted Δ
- ▶ The adversary can *reorder* blocks, always respecting the Δ bound

Network delays



- ▶ Why do we care?
 - ▶ Honest nodes mine on *stale* blocks
 - ▶ **Natural forking** - without adversarial intervention
 - ▶ Waste of honest mining power!

Network delays



- ▶ Ratio of block generation time and network delay is critical for security
- ▶ Bitcoin has low ratio (2%)
- ▶ Ethereum had 12sec block generation time, thus many *orphan* blocks (10-15%)
- ▶ Scalability: we cannot increase arbitrarily the block size or reduce the block generation time without compromising safety!

Network delays

What fraction of honest mining power is wasted due to natural forking?


- ▶ Every time an honest block is created, it takes Δ time to propagate it
- ▶ Expected # blocks mined in Δ is $m_h\Delta$, where m_h is the honest mining rate
- ▶ Only one of these blocks will extend the longest chain
- ▶ Honest miners produce blocks at rate m_h , but only a $\frac{1}{1 + m_h\Delta}$ fraction survives
- ▶ Thus, the rate growth of longest chain is $\frac{m_h}{1 + m_h\Delta}$

Network delays

- ▶ Nakamoto consensus is now secure iff

$$m_a < \frac{m_h}{1 + m_h \Delta}$$


Network delays



Minimum hash power needed by the adversary to succeed in its private attack (security threshold for any successful safety attack).


Selfish mining

Chain quality



- ▶ Chain quality (CQ) can be seen as a measure of *fairness*
- ▶ Fairness *motivates* nodes to participate in the protocol and follow it faithfully
- ▶ Block rewards are awarded accordingly

Chain quality




We showed $CQ \geq 1 - \frac{m_a}{m_h}$, but is this bound tight?

Yes :(


- ▶ Example:
 - ▶ AntPool has 20% of global hash rate
 - ▶ They can gain up to $\frac{20\%}{80\%} = 0.25$ block rewards. Is this possible?

Selfish mining




1. If latest block is H: try to orphan it alone
2. If latest block is A: mine to extend it, announce only once an honest node is produced at the same height.

Selfish mining



1. If latest block is H: try to orphan it alone
2. If latest block is A: mine to extend it, announce only once an honest node is produced at the same height.

Selfish mining



- ▶ Honest nodes and adversary not symmetric
- ▶ Adversary has an advantage by deviating, can reduce fairness
- ▶ Incentives to honestly follow the protocol are not clear
- ▶ Still safety and liveness hold

Selfish mining




Fig. 2: Pool revenue using the Selfish-Mine strategy for different propagation factors γ , compared to the honest Bitcoin mining protocol. Simulation matches the theoretical analysis, and both show that Selfish-Mine results in higher revenues than the honest protocol above a threshold, which depends on γ .




Fig. 3: For a given γ , the threshold α shows the minimum power selfish mining pool that will trump the honest protocol. The current Bitcoin protocol allows $\gamma = 1$, where Selfish-Mine is always superior. Even under unrealistically favorable assumptions, the threshold is never below $1/3$.

Majority is not Enough:
Bitcoin Mining is Vulnerable*

Ittay Eyal and Emin Gün Sirer

Computer Science, Cornell University
ittay.cs.cs.cornell.edu

A sneak peek into the future

More on the economics of blockchains on Lecture 9!



Recap

- ▶ Bitcoin is live but not safe in network partitions
- ▶ Bitcoin is permissionless but not decentralized
- ▶ PoW bypasses FLM: randomization & adversary cannot simulate
- ▶ Nakamoto consensus is secure for minority adversaries assuming 0 network delays in synchrony
- ▶ Network delays affect security by increasing the power of the adversary via natural forking - latency matters
- ▶ Security does not imply Nakamoto consensus is also fair
- ▶ There are attacks (e.g., selfish mining) that give an adversary more block rewards than their fair share

Nakamoto Consensus

Questions?

Bitcoin: A peer-to-peer electronic cash system. Satoshi Nakamoto. 2008.

Majority is not Enough:Bitcoin mining is Vulnerable. Eyal and Shirer.
Communications of ACM, 2018.

November 19th, 2025